THE PM JOB, THEN AND NOW · 2016 VS 2026

From PRDs to Prototypes.

Dennis · Sr PM at Microsoft · @dennisthinks_

1 Start with the problem, not the doc

In 2016 the job opened with a document. The PM gathered requests, wrote the PRD, and handed it to engineering. The doc was the deliverable. In 2026 the doc comes last, after the problem is clear. The work starts earlier: whose problem, how big, and how you would know you fixed it.

A spec written before the problem is clear just makes the wrong thing easier to build.

> EXAMPLE
>
> *"Before the spec: this is for new users who drop in week one. Success is activation up 10 points in 30 days."*

2 Build the prototype before the spec

The 2016 PM described the thing. The 2026 PM builds a rough version of it, puts it in front of a user, and watches. A prototype you can click settles arguments a deck never will.

It does not have to be pretty. It has to be testable. Rough is fine. A user can tap it and show you where it breaks.

> EXAMPLE
>
> *"I built the flow last night. Here it is on my phone. Watch where she hesitates."*

3 Run the eval

AI features changed what checking the work means. An eval is a set of real cases you score the output against, run every time the product changes. It is how you know the new version is better, instead of hoping it is.

Start small. Twenty real user requests, the answer you wanted for each, and a score. That is an eval. Grow it every week.

> EXAMPLE
>
> *"We ran the 40 cases from support tickets. The new version passes 34. Last month's passed 27."*

4 Argue the roadmap with numbers

The 2016 roadmap lived in a deck, and the loudest voice in the review picked the order. The 2026 roadmap still gets argued about. The difference is what gets brought to the argument: usage numbers, the eval scores, the drop-off point with a date on it.

Bring the number. If the number is a guess, say it is a guess and say how you will check it.

> EXAMPLE
>
> *"Onboarding loses about 40% of new users at the invite step. Here is the chart, and here is the fix I want first."*

5 Ship, measure, change one thing

Shipping used to be the finish line. Now launch day is where the checking starts. Ship the smallest version, measure what users do, change the one thing the numbers point at, and ship again.

Practice on this one: take a feature you use every day. Write the problem in one line, sketch the prototype on paper, and list five cases you would score it against. Fifteen minutes.

Score every roadmap

Before the review, ask:

- [ ] Did I name the user and the problem before the feature?
- [ ] Did I show a prototype instead of describing one?
- [ ] Do I have eval cases with scores, not opinions?
- [ ] Does every roadmap item carry a number or a check?
- [ ] Did I ship the smallest version first?
- [ ] Did I change one thing based on what I measured?